

Throughput vs concurrency — vLLM, Qwen2.5-7B-Instruct, FP16

